

A Simple Machine Learning Pipeline


Pipeline ?

sklearn.pipeline.Pipeline


Pipeline as Algorithm Abstraction




Pipeline as Algorithm Abstraction




Pipeline as Algorithm Abstraction




Pipeline as Algorithm Abstraction



Pipeline as Algorithm Abstraction



Pipeline as Algorithm Abstraction



Boston Housing Dataset

<https://archive.ics.uci.edu/ml/datasets/Housing>



<http://bostinno.streetwise.co/2014/12/18/most-expensive-boston-neighborhood-zillow-2014-boston-housing-data/>

```
from sklearn.datasets import load_boston
housing_data = load_boston()
print housing_data.DESCR
```

Boston House Prices dataset

Notes

Data Set Characteristics:

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive

:Median Value (attribute 14) is usually the target

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

```
housing_data.data  
housing_data.target
```



```
import pandas as pd
```

```
df = pd.DataFrame(housing_data.data)
```

```
df.columns = housing_data.feature_names
```

```
df['PRICE'] = housing_data.target
```

```
df.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	PRICE
0	0.00632	18	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2

```
import pandas as pd
```

```
df = pd.DataFrame(housing_data.data)
```

```
df.columns = housing_data.feature_names
```

```
df['PRICE'] = housing_data.target
```

```
df.head()
```

goal



	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	PRICE
0	0.00632	18	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2

Before implementing any algorithm...

Before implementing any algorithm...

... define the metric!


```
import numpy as np
X = df.drop('PRICE', axis=1)
y = df['PRICE']
```

ML-friendly notation
X: feature matrix (506 x 13)
y: target vector (506,)

Cross Validation

	fold 1	fold 2	fold 3
run 1	test		
run 2		test	
run 3			test

```
from sklearn import cross_validation
```

```
def evaluate_model(X, y, algorithm):
```

```
    print 'Mean Square Error'
```

```
    scores = cross_validation.cross_val_score(algorithm, X, y, scoring='mean_squared_error')
```


```
    print -scores
```

```
    print 'Accuracy: %0.2f' % -scores.mean()
```

```
from sklearn import cross_validation
```

```
def evaluate_model(X, y, algorithm):  
    print 'Mean Square Error'  
    scores = cross_validation.cross_val_score(algorithm, X, y, scoring='mean_squared_error')  
    print -scores  
    print 'Accuracy: %0.2f' % -scores.mean()
```

The algorithm – v01



```
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import FunctionTransformer
from sklearn.linear_model import LinearRegression
```


```
def just_RM_column(X):
    RM_col_index = 5
    return X[:, [RM_col_index]]
```

```
pipe = make_pipeline(
    FunctionTransformer(just_RM_column),
    LinearRegression()
)
```

```
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import FunctionTransformer
from sklearn.linear_model import LinearRegression
```

```
def just_RM_column(X):
    RM_col_index = 5
    return X[:, [RM_col_index]]
```

```
pipe = make_pipeline(
    FunctionTransformer(just_RM_column),
    LinearRegression()
)
```






```
evaluate_model(X, y, pipe)
```

Mean Square Error

[43.19492771 41.72813479 46.89293772]

Accuracy: 43.94






The algorithm – v02

```
def add_squared_col(X):  
    return np.hstack((X, X**2))
```

```
pipe = make_pipeline(  
    FunctionTransformer(just_RM_column),  
    FunctionTransformer(add_squared_col),  
    LinearRegression()  
)
```





```
evaluate_model(X, y, pipe)
```

Mean Square Error

[40.31207562 36.75642688 40.75444834]

Accuracy: 39.27






The algorithm – v03

```
from sklearn.tree import DecisionTreeRegressor
```

```
pipe = make_pipeline(  
    FunctionTransformer(just_RM_column),  
    FunctionTransformer(add_squared_col),  
    DecisionTreeRegressor(max_depth=3)  
)
```






```
evaluate_model(X, y, pipe)
```

Mean Square Error


[34.75236642 38.48146015 45.16635916]

Accuracy: 39.47






The algorithm – v04



```
def RM_and_INDUS_cols(X):  
    RM_col_index = 5  
    INDUS_col_index = 2  
    return X[:, [RM_col_index, INDUS_col_index]]
```

```
pipe = make_pipeline(  
    FunctionTransformer(RM_and_INDUS_cols),  
    FunctionTransformer(add_squared_col),  
    LinearRegression()  
)
```



```
evaluate_model(X, y, pipe)
```


Mean Square Error

[32.3420789 31.4260901 35.95835866]

Accuracy: 33.24




see GIF



The algorithm – v05


```
pipe = make_pipeline(  
    FunctionTransformer(add_squared_col),  
    LinearRegression()  
)
```




```
evaluate_model(X, y, pipe)
```

Mean Square Error

[16.7819682 14.599869 18.17785453]

Accuracy: 16.52

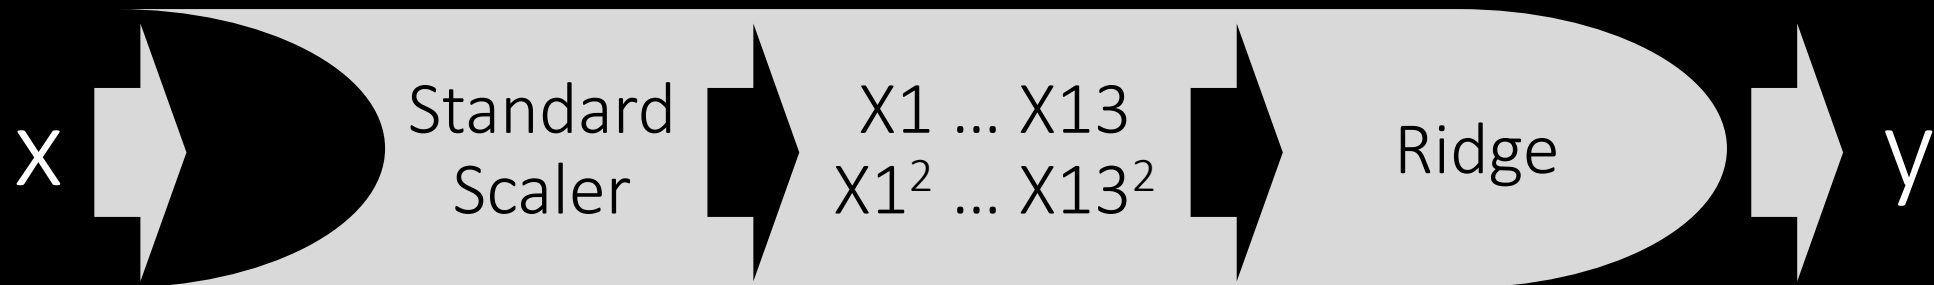


The plot??

The algorithm – v06

```
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Ridge
```

```
pipe = make_pipeline(
    StandardScaler(),
    FunctionTransformer(add_squared_col),
    Ridge(alpha=3)
)
```




```
evaluate_model(X, y, pipe)
```

Mean Square Error


[16.4292824 14.50522561 18.27167008]

Accuracy: 16.40



In conclusion...

Pipeline as Algorithm Abstraction



Questions?

Want more?
Data Science Milan

tw: @MrSantoni
marcosantoni.com